


VIT UNIVERSITY, ANDHRA PRADESH
School of CSE
CSE3008 - Introduction to Machine Learning
Lab Experiment-10
(**Multiclass Classification using Support Vector Machine,
Support vector regression**)
Faculty-**Dr. B. SRINIVASA RAO**

Name-Neeraj Guntuku
R.No-18MIS7071
Slot-L55+L56

Date- 10 April 2021

Multiclass Classification using Support Vector Machine



▾ Training Support Vector Machines for Multiclass Classification

```
[1]
import numpy as np
import pylab as pl
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn.utils import shuffle
from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import cross_val_score, GridSearchCV

import os
print(os.listdir())

['.config', 'train.csv', 'test.csv', '.ipynb_checkpoints', 'sample_data']
```

▼ Load the Train and Test set

```
[2] train = shuffle(pd.read_csv("train.csv"))
     test = shuffle(pd.read_csv("test.csv"))
```

▼ Check for missing values in the dataset

```
[3] print("Any missing sample in training set:",train.isnull().values.any())
     print("Any missing sample in test set:",test.isnull().values.any(), "\n")
```

```
Any missing sample in training set: False
Any missing sample in test set: False
```

▼ Frequency Distribution of the Outcome

```
[4] #Frequency distribution of classes"
     train_outcome = pd.crosstab(index=train["Activity"], # Make a crosstab
                                columns="count")        # Name the count column

     train_outcome
```

	col_0	count
Activity		
LAYING		1407
SITTING		1286
STANDING		1374
WALKING		1226
WALKING_DOWNSTAIRS		986
WALKING_UPSTAIRS		1073

Visualizing Outcome Distribution

```
[5] # Visualizing Outcome Distribution
temp = train["Activity"].value_counts()
df = pd.DataFrame({'labels': temp.index,
                   'values': temp.values
                   })

#df.plot(kind='pie',labels='labels',values='values', title='Activity Distribution',subplots= "True")

labels = df['labels']
sizes = df['values']
colors = ['yellowgreen', 'gold', 'lightskyblue', 'lightcoral','cyan','lightpink']
patches, texts = plt.pie(sizes, colors=colors, shadow=True, startangle=90, pctdistance=1.1, labeldistan
plt.legend(patches, labels, loc="best")
plt.axis('equal')
plt.tight_layout()
plt.show()
```



Normalize the Predictor(Feature Set) for SVM training

```
[6] # Separating Predictors and Outcome values from train and test sets
X_train = pd.DataFrame(train.drop(['Activity', 'subject'],axis=1))
Y_train_label = train.Activity.values.astype(str)
X_test = pd.DataFrame(test.drop(['Activity', 'subject'],axis=1))
Y_test_label = test.Activity.values.astype(str)

# Dimension of Train and Test set
print("Dimension of Train set",X_train.shape)
print("Dimension of Test set",X_test.shape,"\n")

# Transforming non numerical labels into numerical labels
from sklearn import preprocessing
encoder = preprocessing.LabelEncoder()

# encoding train labels
encoder.fit(Y_train_label)
Y_train = encoder.transform(Y_train_label)

# encoding test labels
encoder.fit(Y_test_label)
Y_test = encoder.transform(Y_test_label)

#Total Number of Continous and Categorical features in the training set
num_cols = X_train._get_numeric_data().columns
print("Number of numeric features:",num_cols.size)
#list(set(X_train.columns) - set(num_cols))
```

```

names_of_predictors = list(X_train.columns.values)

# Scaling the Train and Test feature set
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

Dimension of Train set (7352, 561)
Dimension of Test set (2947, 561)

Number of numeric features: 561

▼ Hyperparameter tuning using grid search and cross validation

```

[7] #Libraries to Build Ensemble Model : Random Forest Classifier
# Create the parameter grid based on the results of random search
params_grid = [{'kernel': ['rbf'], 'gamma': [1e-3, 1e-4],
                        'C': [1, 10, 100, 1000]},
                {'kernel': ['linear'], 'C': [1, 10, 100, 1000]}]

```

▼ Training SVM model using radial kernel

```

[8] # Performing CV to tune parameters for best SVM fit
svm_model = GridSearchCV(SVC(), params_grid, cv=5)
svm_model.fit(X_train_scaled, Y_train)

```

```

GridSearchCV(cv=5, error_score=nan,
             estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                           class_weight=None, coef0=0.0,
                           decision_function_shape='ovr', degree=3,
                           gamma='scale', kernel='rbf', max_iter=-1,
                           probability=False, random_state=None, shrinking=True,
                           tol=0.001, verbose=False),
             iid='deprecated', n_jobs=None,
             param_grid=[{'C': [1, 10, 100, 1000], 'gamma': [0.001, 0.0001],
                           'kernel': ['rbf']},
                          {'C': [1, 10, 100, 1000], 'kernel': ['linear']}],
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=0)

```

▼ Confusion Matrix and Accuracy Score

```
[9] # View the accuracy score
print('Best score for training data:', svm_model.best_score_,"\n")

# View the best parameters for the model found using grid search
print('Best C:',svm_model.best_estimator_.C,"\n")
print('Best Kernel:',svm_model.best_estimator_.kernel,"\n")
print('Best Gamma:',svm_model.best_estimator_.gamma,"\n")

final_model = svm_model.best_estimator_
Y_pred = final_model.predict(X_test_scaled)
Y_pred_label = list(encoder.inverse_transform(Y_pred))
```

Best score for training data: 0.9865340344159417

Best C: 1000

Best Kernel: rbf

Best Gamma: 0.001

```
[10] # Making the Confusion Matrix
#print(pd.crosstab(Y_test_label, Y_pred_label, rownames=['Actual Activity'], colnames=['Predicted Activ
print(confusion_matrix(Y_test_label,Y_pred_label))
print("\n")
print(classification_report(Y_test_label,Y_pred_label))
```

```
print("Training set score for SVM: %f" % final_model.score(X_train_scaled , Y_train))
print("Testing set score for SVM: %f" % final_model.score(X_test_scaled , Y_test ))
```

svm_model.score

```
[[537  0  0  0  0  0]
 [ 3 439 48  0  0  1]
 [ 0 11 521  0  0  0]
 [ 0  0  0 486  4  6]
 [ 0  0  0  6 389 25]
 [ 0  0  0 15  2 454]]
```

```
[[537  0  0  0  0  0]
 [  3 439 48  0  0  1]
 [  0 11 521  0  0  0]
 [  0  0  0 486  4  6]
 [  0  0  0  6 389 25]
 [  0  0  0 15  2 454]]
```

	precision	recall	f1-score	support
LAYING	0.99	1.00	1.00	537
SITTING	0.98	0.89	0.93	491
STANDING	0.92	0.98	0.95	532
WALKING	0.96	0.98	0.97	496
WALKING_DOWNSTAIRS	0.98	0.93	0.95	420
WALKING_UPSTAIRS	0.93	0.96	0.95	471
accuracy			0.96	2947
macro avg	0.96	0.96	0.96	2947
weighted avg	0.96	0.96	0.96	2947

Training set score for SVM: 1.000000

Testing set score for SVM: 0.958941

```
<bound method BaseSearchCV.score of GridSearchCV(cv=5, error_score=nan,
          estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                        class_weight=None, coef0=0.0,
                        decision_function_shape='ovr', degree=3,
                        gamma='scale', kernel='rbf', max_iter=-1,
                        probability=False, random_state=None, shrinking=True,
                        tol=0.001, verbose=False),
          iid='deprecated', n_jobs=None,
          param_grid=[{'C': [1, 10, 100, 1000], 'gamma': [0.001, 0.0001],
                       'kernel': ['rbf']},
                      {'C': [1, 10, 100, 1000], 'kernel': ['linear']}],
          pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
          scoring=None, verbose=0)>
```

Support vector regression

Support Vector Regression

Importing the libraries

```
[1] import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

Importing the dataset

```
[2] dataset = pd.read_csv('SampleData.csv')
X = dataset.iloc[:, 0].values
y = dataset.iloc[:, 1].values
y = np.array(y).reshape(-1,1)
dataset.head(5)
```

	Hours of Study	Marks
0	32.502345	31.707006
1	53.426804	68.777596
2	61.530358	62.562382
3	47.475640	71.546632
4	59.813208	87.230925

Feature Scaling

```
[3] from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
sc_y = StandardScaler()
X = sc_X.fit_transform(X.reshape(-1,1))
y = sc_y.fit_transform(y.reshape(-1,1))
```

Splitting the dataset into the Training set and Test set

```
[4] from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
```

▼ Training the Support Vector Regression model on the Training set

```
[5] from sklearn.svm import SVR
regressor = SVR(kernel = 'rbf')
regressor.fit(X_train.reshape(-1,1), y_train.reshape(-1,1))

/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:760: DataConversionWarning: A
y = column_or_1d(y, warn=True)
SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1, gamma='scale',
kernel='rbf', max_iter=-1, shrinking=True, tol=0.001, verbose=False)
```

▼ Predicting the Results

```
[6] y_pred = regressor.predict(X_test)
y_pred = sc_y.inverse_transform(y_pred)
y_pred

array([ 78.86474326,  93.07233808,  82.04137655,  79.79709186,  76.72265591,
        90.54612743,  60.38192467,  58.55099588,  71.12798091,  62.11580563,
        69.89505864,  59.0668753 ,  61.31339316,  84.32726724,  85.42956343,
        59.33656194,  74.69547341,  76.54528267,  53.41532363,  68.30603506])
```

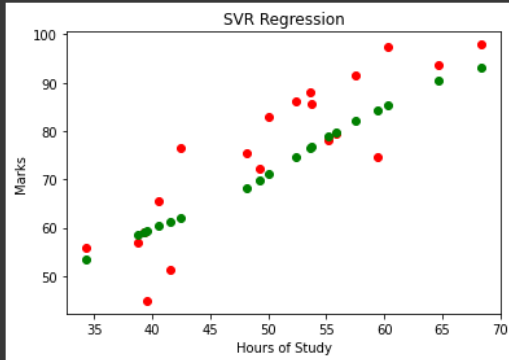
▼ Comparing the Real Values with Predicted Values

```
[7] df = pd.DataFrame({'Real Values':sc_y.inverse_transform(y_test.reshape(-1)), 'Predicted Values':y_pred})
df
```


	Real Values	Predicted Values
0	78.211518	78.864743
1	97.919821	93.072338
2	91.486778	82.041377
3	79.550437	79.797092
4	85.668203	76.722656
5	93.576119	90.546127
6	65.562301	60.381925
7	56.877213	58.550996
8	82.905981	71.127981
9	76.617341	62.115806
10	72.111832	69.895059
11	59.171489	59.066875
12	51.391744	61.313393
13	74.765564	84.327267
14	97.379897	85.429563
15	44.888491	53.415323

▸ Visualising the SVR Results

```
[8] # Visualising the SVR results (for higher resolution and smoother curve)
X_grid = np.arange(min(X), max(X), 0.1)
X_grid = X_grid.reshape((len(X_grid), 1))
plt.scatter(sc_X.inverse_transform(X_test), sc_y.inverse_transform(y_test.reshape(-1)), color = 'red')
plt.scatter(sc_X.inverse_transform(X_test), y_pred, color = 'green')
plt.title('SVR Regression')
plt.xlabel('Hours of Study')
plt.ylabel('Marks')
plt.show()
```



```
[9] plt.plot(X_grid, regressor.predict(X_grid), color = 'black')
plt.title('SVR Regression')
plt.xlabel('Hours of Study')
plt.ylabel('Marks')
plt.show()
```



..

sample_data

SampleData.csv
